

# Character-based Embeddings of Words with Recurrent Nets

**Simon Grätzer**

**Seminar Selected Topics in Human Language Technology and Pattern Recognition – MedBV16 June 10, 2016**

**Human Language Technology and Pattern Recognition  
Lehrstuhl für Informatik 6  
Computer Science Department  
RWTH Aachen University, Germany**

# Outline

- 1. Introduction**
- 2. Word Embeddings**
- 3. Generating Word Embeddings**


# Literature

- ▶
- ▶ **[?] The tensorflow doumentation by Google Inc.**

# Introduction

- ▶ Word embeddings are real valued vector representations for words.
- ▶ In this talk I will present a new idea to generate these representations.
  - ▷ Using recurrent neural networks (LSTM)
  - ▷ Using individual representations of the characters as inputs.
- ▶ The resulting model can be used to improve some tasks, such as language modeling or part-of-speech tagging

AUDIO



Audio Spectrogram

DENSE

IMAGES





Image pixels

DENSE

TEXT


Word, context, or  
document vectors

SPARSE


Different input datasets compared to word data [?].

**The underlying problem is the sparsity of ordinary word representations.**


# Word Embeddings



Male-Female




Verb tense



Country-Capital

- ▶ Words which share semantic meaning tend to occur in the same contexts (Distributional Hypothesis [?])
- ▶ A model can learn relationships between words and represent them.
- ▶ Words with a similar meaning should be mapped to nearby points in the same vector space.
- ▶ This captures the intuition that words may be similar along a variety of ways.


# Simple Word Embeddings in a Language Model



**Classical natural language model**

- Language modelling estimates  $p(w_1, \dots, w_m) = \prod_{i=1}^m p(w_i | w_1, \dots, w_{i-1})$
- The context is approximated with the previous  $n - 1$  words (n-grams)
- $p(w_1, \dots, w_m) = \prod_{i=1}^m p(w_i | w_1, \dots, w_{i-1}) \approx \prod_{i=1}^m p(w_i | w_{i-(n-1)}, \dots, w_{i-1})$


# Simple Word Embeddings in a Language Model



**Classical natural language model**

- ▶ Word Embeddings are trained end to end, as part of the language model
- ▶ The word embeddings is essentially the output of the first hidden layer for a word.
- ▶ The embeddings are stored in a lookup table  $P \in \mathbb{R}^{|V| \times d}$ . An embeddings is calculated as  $e_{w_i}^W = P * 1_{w_i}$

# Repetition: Softmax-Layer




Input  
(Features II)      Softmax  
                         classifier

$$p_k = \sigma(\mathbf{z})_k = \frac{e^{z_k}}{\sum_{k=1}^{|V|} e^{z_k}}$$


- ▶ The input vector  $\mathbf{z}$  is computed over the vocabulary:  $z_k \forall k \in V$ .
- ▶ The result of the outputs can be interpreted as posterior probabilities.
- ▶ Probability given the context:  $p_k = p(w_i = k | w_{i-n+1}^{i-1})$




# Advanced Word Embeddings: Word2vec



Male-Female



Verb tense



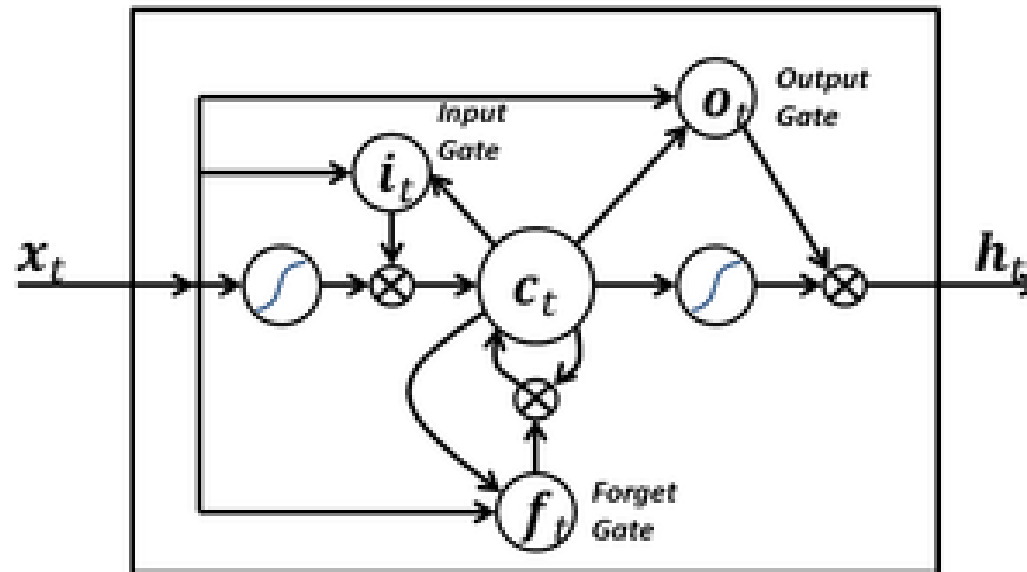
Country-Capital

- ▶ The Skip-Gram based model from Mikolov et.al. was developed at Google.
- ▶ They use n-gram's but invert the model: 
$$\sum_{-k \leq j-1, j \leq k} \log P(w_{t+j} | w_t)$$
- ▶  $v(\text{king}) - v(\text{male}) + v(\text{female}) \approx v(\text{queen})$
- ▶ Resulting lookup table of embeddings can be reused for other tasks.

# Drawbacks of these Embeddings

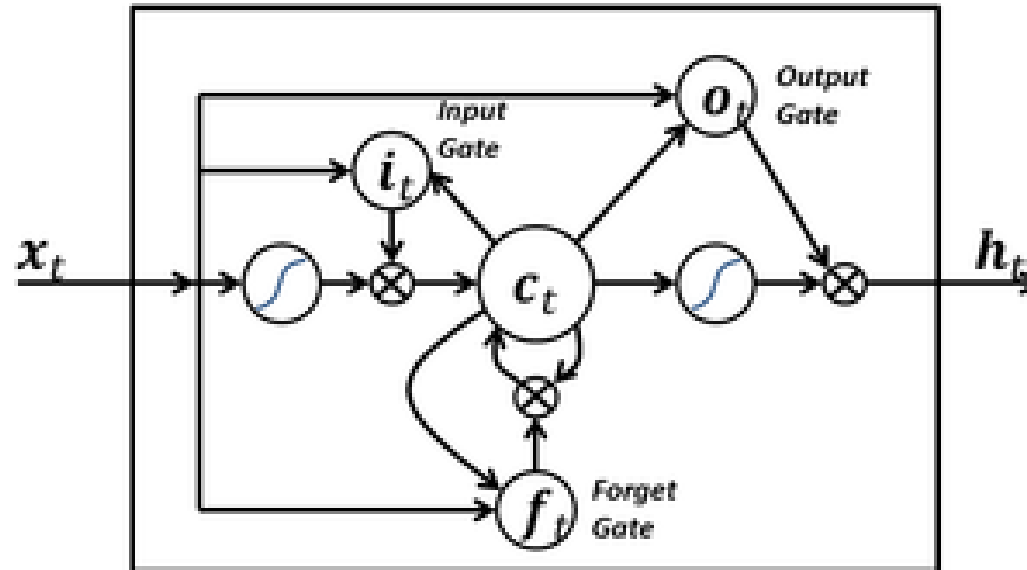
- ▶ **Each word embedding vector is completely independent.**
  - ▷ **The model captures similar linear correspondences between words embeddings i.e. *cat* and *apple* compared to *cats* and *apples***
  - ▷ **It doesn't capture that the added *s* is responsible for this transformation.**
  - ▷ **A word lookup table cannot generate representations for an unknown word.**
  - ▷ **Even if it's just the plural form of a known word.**
- ▶ **For a large vocabulary it becomes impractical to actually store all word embeddings in a table.**

# Repetition: Long-Short Term Memory



- ▶ Designed to "remember" inputs over long distances and "forget" them when necessary
- ▶ Works well for time series data
- ▶ Gate  $i_t$  to determine when to learn an input value
- ▶ Gate  $f_t$  to determine if it should continue to remember or forget the currently stored value
- ▶ Gate  $o_t$  to determine whether it should output the value.

# Repetition: Long-Short Term Memory

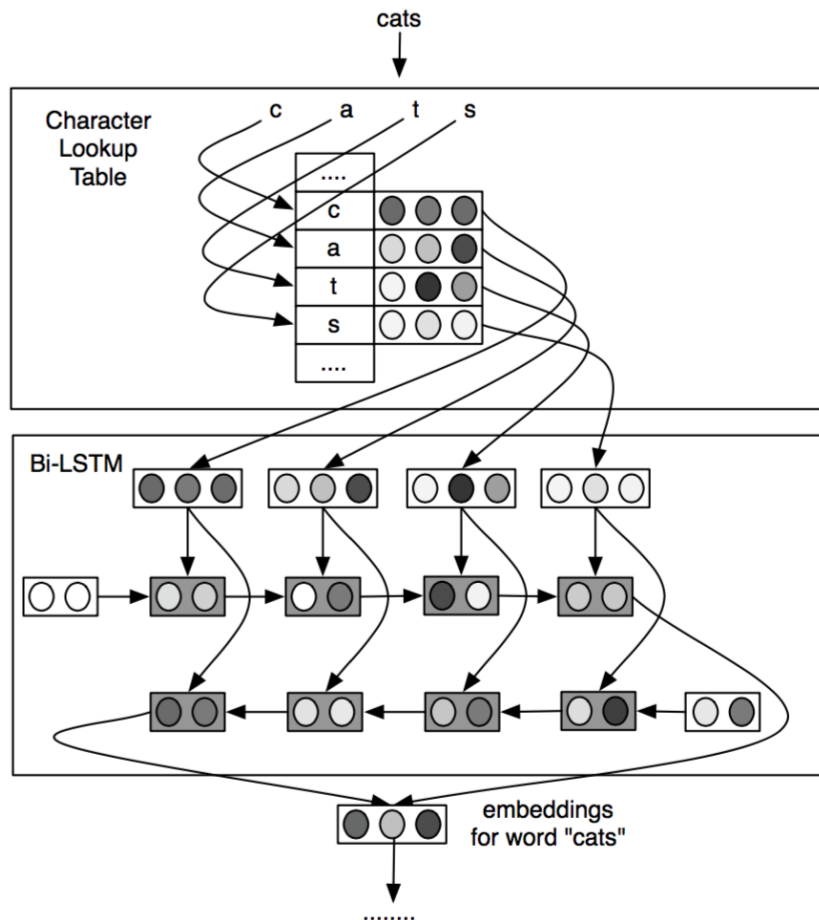


Given the input vectors  $x_1, \dots, x_m$  a LSTM computes the output sequence  $h_1, \dots, h_m$

$$\begin{aligned}
 i_t &= \sigma(W_{ix} * x_t + W_{ih} * h_{t-1} + W_{ic} * c_{t-1} + b_i) \\
 f_t &= \sigma(W_{fx} * x_t + W_{fh} * h_{t-1} + W_{fc} * c_{t-1} + b_f) \\
 o_t &= \sigma(W_{ox} * x_t + W_{oh} * h_{t-1} + W_{oc} * c_t + b_o) \\
 c_t &= f_t \odot c_{t-1} + i_t \odot \tanh(W_{cx} * x_t + W_{ch} * h_{t-1} + b_c) \\
 h_t &= o_t \odot \tanh(c_t)
 \end{aligned} \tag{1}$$

**LSTM's avoid the vanishing gradient problem, because the activation function in  $c_t$  is the identity function.**

# Character-based Word-Embeddings (C2W)



Character lookup table on top,  
bidirectional LSTM on the  
bottom


## Overview:

- ▶ A word with length  $m$  is composed of characters  $c_1, \dots, c_m$ .
- ▶ Decompose each word into a sequence of character embeddings  $e_{c_1}^C, \dots, e_{c_m}^C$  from the alphabet  $C$ .
- ▶ The sequence is fed to two LSTM units (Bidirectional LSTM), forwards and backwards.
- ▶ In the end the result is combined by an output layer.

# C2W-Model: Character-Lookup Table

- ▶ Table of  $d_C$  parameters  $P_C \in \mathbb{R}^{d_C \times |C|}$  for each character from a predefined character alphabet  $C$ .
- ▶ Each Input character is transformed into a  $d_C$ -dimensional feature vector  $e_{c_j}^C$ .
- ▶ We define the projection of characters as  $e_{c_j}^C = P_C * 1_{c_j}$ .
- ▶ Similar to the previous projection layer for each word.

# C2W-Model: Bidirectional LSTM Layer



- ▶ Present every input sequence forwards and backwards to two separate recurrent neural networks
- ▶ Both RNN's are connected to the same output layer.
- ▶ The network has simultaneous access to all inputs before and after the current one.
- ▶ No need for fixed window sizes for the input, the net decides how much context to use.



- Yields the forward state sequence  $s_0^f, \dots, s_m^f$  and backward state sequence  $s_m^b, \dots, s_0^b$ .

## C2W-Model: Output Layer

- ▶ Table of  $d_C$  parameters  $P_C \in \mathbb{R}^{d_C \times |C|}$  for each character from a predefined character alphabet  $C$ .
- ▶ Each Input character is transformed into a  $d_C$ -dimensional feature vector  $e_{c_j}^C$ .
- ▶ We define the projection of characters as  $e_{c_j}^C = P_C * 1_{c_j}$ .
- ▶ Similar to the previous projection layer for each word.

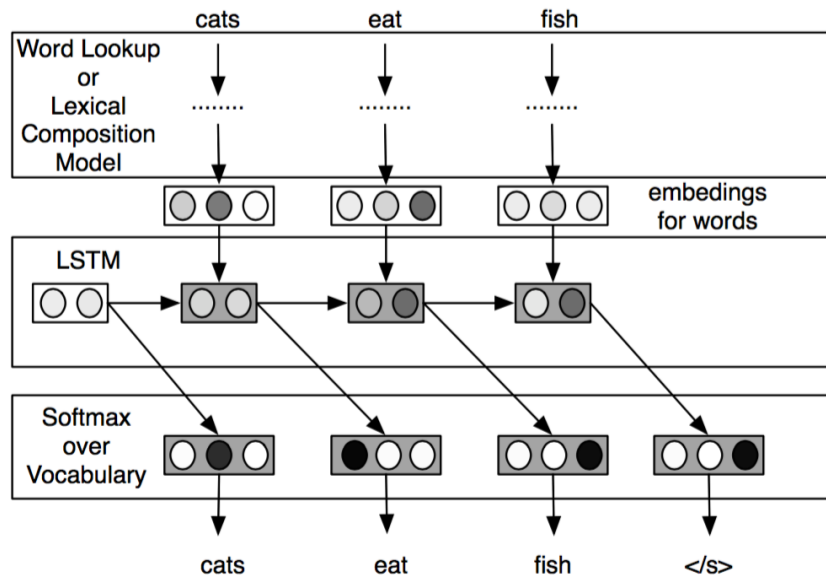
# Character-based Word-Embeddings: Advantages

- ▶ Simply breaks-up words into simple atomic units.
- ▶ Characters are the simplest atomic unit of words.
- ▶ Alternative: Use Morphemes as atomic units
  - ▷ A morpheme is the smallest grammatical unit of a language
  - ▷ e.g. "Unbreakable" comprises three morphemes: un-, -break-, and -able.
  - ▷ Would require a morphological analyser

# Output-Layer

- ▶ Combines the last states of the forward sequence  $s_m^f$  and the backwards sequence  $s_0^b$
- ▶  $e_w^C = D^f s_m^f + D^b s_0^b + b_d$
- ▶ The variables  $D^f, D^b, b_d$  are the weights which determine how the states are combined.
- ▶ Automatically learns how much each context is used.

# Application: Language Modeling



- ▶ Uses the word embeddings from the C2W model combined with a LSTM unit.
- ▶ Every time we input a new word  $w_i$  from the sequence the model yields the LSTM state  $s_i$ .
- ▶ In the end a softmax layer is used to compute the likelihood  $p(w_i = k | w_{i-n+1}^{i-1})$

# Application: Language Modeling - Evaluation

- ▶ **TODO include the evaluation tables**

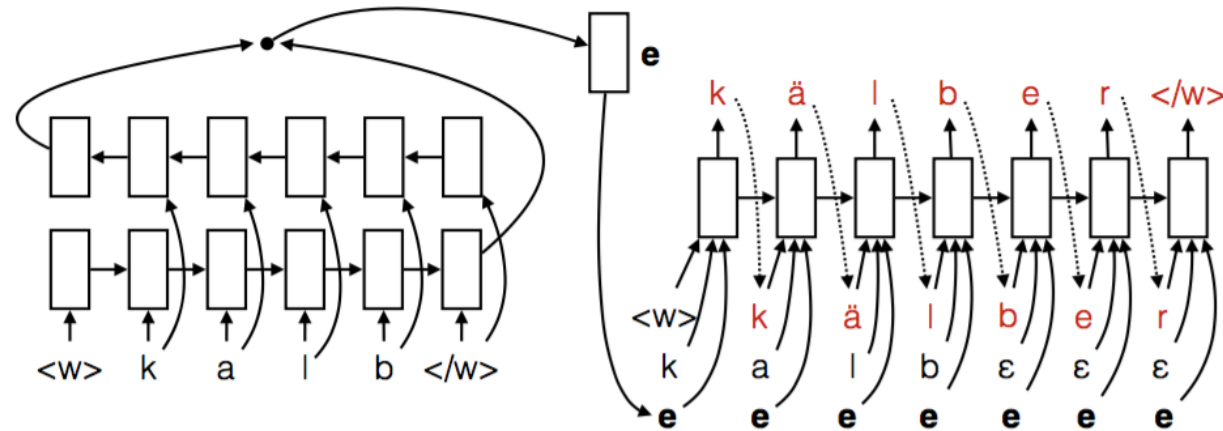
# Applications: Morphological Inflection Generation

	singular	plural
nominative	Kalb	Kälber
accusative	Kalb	Kälber
dative	Kalb	Kälbern
genitive	Kalbes	Kälber

**Example of an inflection table for the word "Kalb"**

- ▶ **Perform morphological transformations, as discussed in Faruqui et al. [?]**
- ▶ **The transformations are very common in languages like turkish or german.**
- ▶ **Basic idea is to use a neuronal encoder - decoder architecture.**
- ▶ **The encoder mirrors the C2W model.**

# Application: Morphological Inflection Generation



- First: Encoder is virtually identically to C2W model and generates a word embedding  $e_w$ .
- The decoder is just an LSTM unit which receives the following inputs each timestep:
  1. The word embedding  $e_w$  from the encoder.
  2. Current character of the original word  $c_j$
  3. Previous output of the model
- Once the input word ends, the  $\epsilon$  character is used instead.



# Summary

- ▶ **Calculating the word embeddings is cheaper than storing them in huge lookup-tables**
- ▶ **Performance is comparable to other methods**
- ▶ **Lexical features can be learned automatically**
- ▶ **Redundancies in lookup-tables are avoided**
- ▶ **Models scale better with larger vocabularies.**

# Backup: Perplexity

- ▶ **Measure of how well a probability distribution predicts sample data.**
- ▶ **Can be interpreted as the number of choices per word position.**
- ▶ **Defined as  $2^{H(p)} = 2^{-\sum_x p(x) \log_2 p(x)}$**
- ▶ **To minimize the perplexity value means to have a better fitting language model.**

# Backup: Out of Vocabulary Token



# Backup: Out of Vocabulary Token



