

# Character-based Embeddings of Words with Recurrent Nets

**Simon Grätzer**

**Seminar Selected Topics in Human Language Technology and Pattern Recognition June 20, 2016**

**Human Language Technology and Pattern Recognition  
Lehrstuhl für Informatik 6  
Computer Science Department  
RWTH Aachen University, Germany**

# Outline

## 1. Introduction

## 2. Word Embeddings

### (a) Goals

### (b) Continuous Space Language Model

### (c) Shortcomings of Word Lookup Tables

## 3. Character-based Word-Embeddings

## 4. Experimentation

### (a) Language Modeling

### (b) Part-Of-Speech Tagging

### (c) Morphological Inflection Generation

# Literature

**Wang Ling et al 2015** : Finding Function in Form: Compositional Character Models for Open Vocabulary Word Representation. CoRR Volume abs/1508.02096, 2015

- ▶ Introduces a model for constructing vector representations of words by composing characters using bidirectional LSTMs.

**Faruqui et al 2015** : Morphological Inflection Generation Using Character Sequence to Sequence Learning. CoRR Volume abs/1512.06110, 2015

- ▶ Approach to generate inflected versions of words by modelling the process as a character sequence to sequence learning problem.

**Mikolov et al 2013** : Distributed Representations of Words and Phrases and their Compositionality. CoRR Volume abs/1310.4546, 2013

- ▶ Improvements on the Skip-gram model used to generate word embeddings.

# Literature

**Schwenk 2007** : Continuous Space Language Models. Journal Computer Speech and Language. Volume 21 Issue 3, July, 2007 Pages 492-518

- ▶ Describes the use of a neural network language model for large vocabulary continuous speech recognition.

**Peirsman 2015** : Visualizing Word Embeddings with t-SNE. Online 2015

- ▶ Creating useful low-dimensional visualizations for high-dimensional datasets

**TensorFlow 2016** : Vector Representations of Words, From the TensorFlow documentation


- ▶ Implementation of word2vec model of [Mikolov et al 2013] with the TensorFlow framework.

# Introduction

- ▶ **Word embeddings are real valued vector representations for words.**
- ▶ **This talk is about generating word embeddings and their applications.**
- ▶ **Specifically:**
  - ▷ **Generating word embeddings by composing their individual character representations**
  - ▷ **Using Long short-term memory to capture complex relationships between words.**
- ▶ **The resulting model can be used for many tasks, such as language modeling or part-of-speech tagging**
- ▶ **Reduces the need for manual feature engineering**
- ▶ **Can improve performance for many tasks**

# Introduction

AUDIO



Audio Spectrogram

DENSE

IMAGES




Image pixels

DENSE

TEXT

0	0	0	0.2	0	0.7	0	0	0	...	...
---	---	---	-----	---	-----	---	---	---	-----	-----


Word, context, or  
document vectors

SPARSE

[TensorFlow 2016]

- ▶ Natural language processing systems can treat words as atomic symbols, encoded as indexes.
- ▶ E.g. 'apple' might become index *ld123* and 'orange' becomes index *ld124*
- ▶ This way of encoding is very sparse, and arbitrary.
- ▶ No representation of the relationship between these word types (such as both are fruit, ...).

# Word Embeddings



Embeddings visualized [Peirson 2015]

- ▶ A statistical model should be able to learn relationships between word types.
- ▶ Transform a word type  $w$  and turn it into an embedding  $e_w = v(w) \in \mathbb{R}^d$
- ▶ Words with a similar meaning should be mapped to (geometrically) nearby points in the vector space.
- ▶ The dimension  $d$  should be low compared to the vocabulary size.
- ▶ Capture the intuition that words may be similar along different dimensions.

# Word Embeddings

**To this end there are some assumptions we make:**

- ▶ **Words which share semantic meaning tend to occur in the same contexts (Distributional Hypothesis)**
- ▶ **The composition of words themselves can sometimes hint at similarities (e.g. 'apple' vs 'apples')**




# General Approach for Word Embeddings

**Generating word embeddings within a NNLM to better estimate  $p(w_i | w_1, \dots, w_{i-1})$**

- 1. Associate each word  $w$  in the vocabulary  $V$  with a word embedding  $e_w$**
- 2. Express the joint probability function for the word sequence  $w_1, \dots, w_{i-1}$  in terms of these embeddings.**
- 3. Simultaneously learn the word embeddings and the parameters of the probability function.**


# Continuous Space Language Model




- ▶ The context for  $w_i$  is approximated with n-grams.
- ▶ Input word types are encoded as one-hot vectors.
- ▶ First Hidden Layer: Lookup table  $P \in \mathbb{R}^{d \times |V|}$ , projects the input into a continuous vector space.
- ▶ The word embedding  $e_w$  is the output of layer  $P$ .
- ▶ Second Hidden Layer: Estimates the joint probability of the word sequence
- ▶ The Softmax-Layer projects this up to vocabulary size  $|V|$

[Schwenk 2007]


# Resulting Embeddings



Male-Female



Verb tense



Country-Capital

[TensorFlow 2016]

- ▶ This setup can be used to generate a large table of word embeddings.
- ▶ Lookup table can be reused for other tasks (if there is not enough training data).
- ▶ E.g. [Mikolov et al 2013] have created a model which can encode complex patterns:
- ▶  $v(\text{king}) - v(\text{male}) + v(\text{female}) \approx v(\text{queen})$ .

# Shortcomings

- ▶ A model with a lookup table treats each word embedding as independent from each other.
  - ▷ The model captures similar linear correspondences between word embeddings.
  - ▷ E.g. *cat* and *apple* compared to *cats* and *apples*.
  - ▷ It doesn't capture that the added *s* is responsible for this transformation.
  - ▷ The model doesn't examine lexical similarities between words.
  - ▷ It doesn't capture morphological word transformations: e.g. *test* vs. *testing*
- ▶ A word lookup table cannot easily deal with unknown words.
- ▶ The lookup table contains at least  $|V| \times d$  parameters. This can require large amount of memory for tasks with large vocabularies.

# Possible Solutions

Some requirements should be satisfied by a better model:

- ▶ The model should capture\* orthographic similarities between words e.g. *test* vs. *testing* (Compositional effects).
- ▶ The model must still capture functionally similar words, with no orthographic similarities e.g. *rich* vs. *affluent*.
- ▶ However not all similar spelled words have similar meanings e.g. *butter* vs. *batter*.
- ▶ The resulting model should be able to replace the previous projection layer.
- ▶ Idea: Break down words into smaller atomic units and try to *compose* them into word embeddings.

\*) In terms of geometric locality

# Possible Solution 1: Morpheme-based Embeddings

- ▶ **Morphemes are the smallest defined grammatical unit of a language.**
- ▶ **E.g. "Unbreakable" comprises three morphemes:**
  - 1. un-**
  - 2. -break-**
  - 3. -able**
- ▶ **Use morphemes as input and compose them into the word embeddings.**
- ▶ **Requires a morphological analyser: Extra processing step.**
- ▶ **Each target language requires an extra morphological analyser.**
- ▶ **The quality of the model depends on the analyser.**

# Possible Solution 2: Character-based Embeddings

- ▶ **Break down words into characters**
- ▶ **Represent characters as real valued feature vectors (character embeddings).**
- ▶ **Feed them to an recurrent neural net, which "remembers" each character.**
- ▶ **Learn character embeddings simultaneously with other model parameters.**

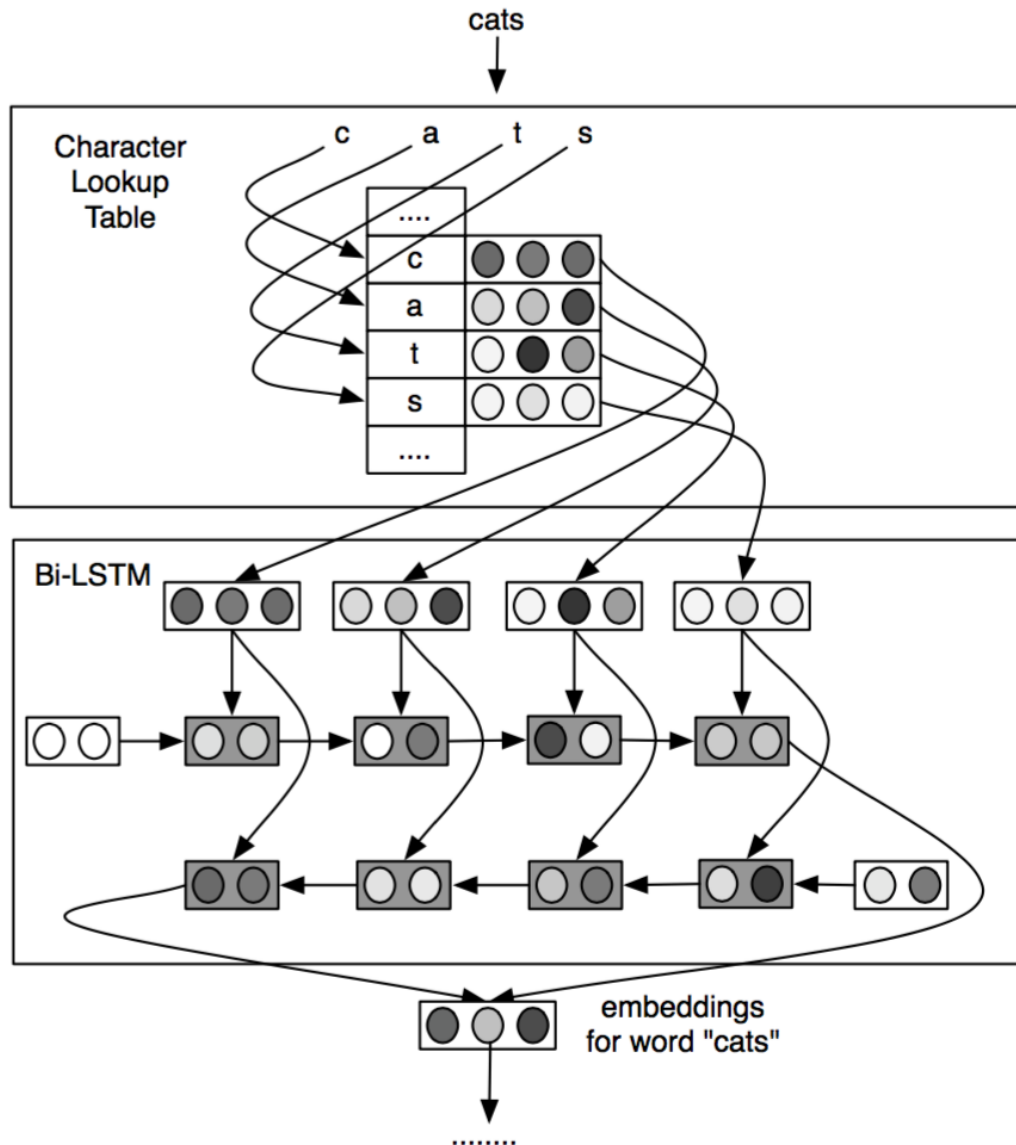
# Character-based Word-Embeddings

The "Compositional Character to Word" (C2W) model:

1. A word  $w$  with length  $m$  is decomposed into characters  $c_1, \dots, c_m$  from the alphabet  $C$ .
2. Transform characters into a sequence of character embeddings  $e_{c_1}, \dots, e_{c_m}$ .
3. The sequence is "read" one-by-one forwards as well as backwards by two LSTMs.
4. During "reading" the sequence is composed into the forward state  $s_m^f$  and backward state  $s_1^b$ .
5. The two states are recomposed to form the word-embedding  $e_w$ .



# Compositional Character to Word Model




[Wang Ling et al 2015]

The C2W model is visualized for input "cats":

- ▶ Squared boxes represent vectors of neuron activations.
- ▶ Shaded boxes indicate a nonlinear output.
- ▶ The two actual LSTM units are displayed unfolded.


# C2W-Model: Layers



[Wang Ling et al 2015]

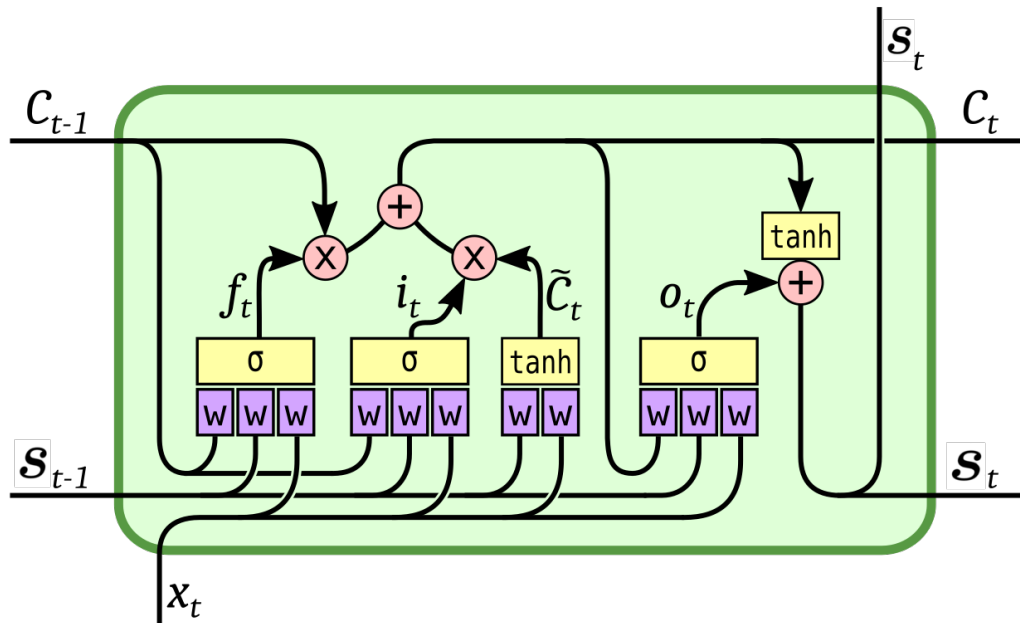
1. Table of character embeddings.
2. Bidirectional-LSTM layer, processing forward- and backward-sequences of character embeddings.
3. The combining layer, to merge the two outputs from the Bi-LSTM.

# C2W-Model: Character-Lookup Table



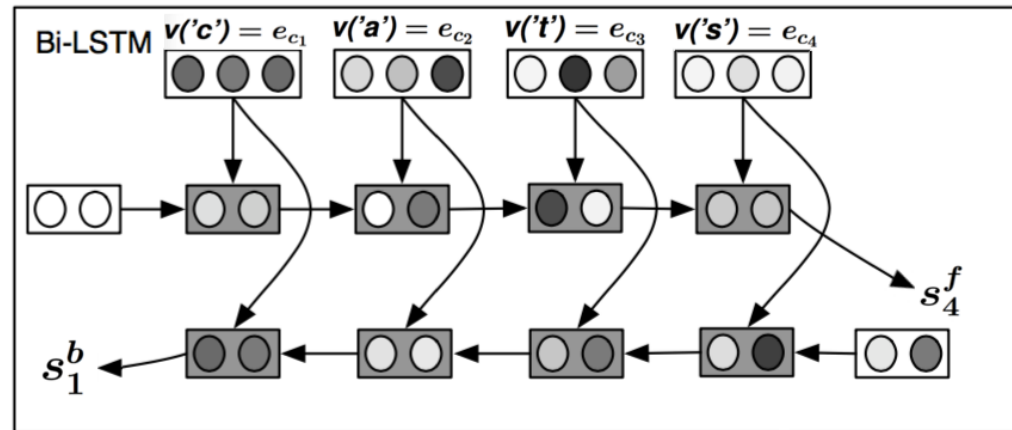
- ▶ Table of  $d_C$  parameters  $P_C \in \mathbb{R}^{d_C \times |C|}$ .
- ▶ Each Input character  $c$  is transformed into a  $d_C$ -dimensional feature vector  $e_c$ .
- ▶ The dimension  $d_C$  becomes a hyperparameter of the model.
- ▶ Basically similar to the previous projection layer for words.

# Reminder: Long Short-Term Memory



- ▶ Designed to "remember" inputs over arbitrary distances and "forget" them when necessary
- ▶ Gate  $i_t$  to determine when to learn an input value
- ▶ Gate  $f_t$  to determine if it should continue to remember or forget the currently stored value
- ▶ Gate  $o_t$  to determine whether it should output the value.
- ▶ Additionally and bias values not explicitly displayed here.
- ▶ The dimension  $d_{CS}$  of the LSTM state becomes another hyperparameter.

# C2W-Model: Bidirectional LSTM Layer

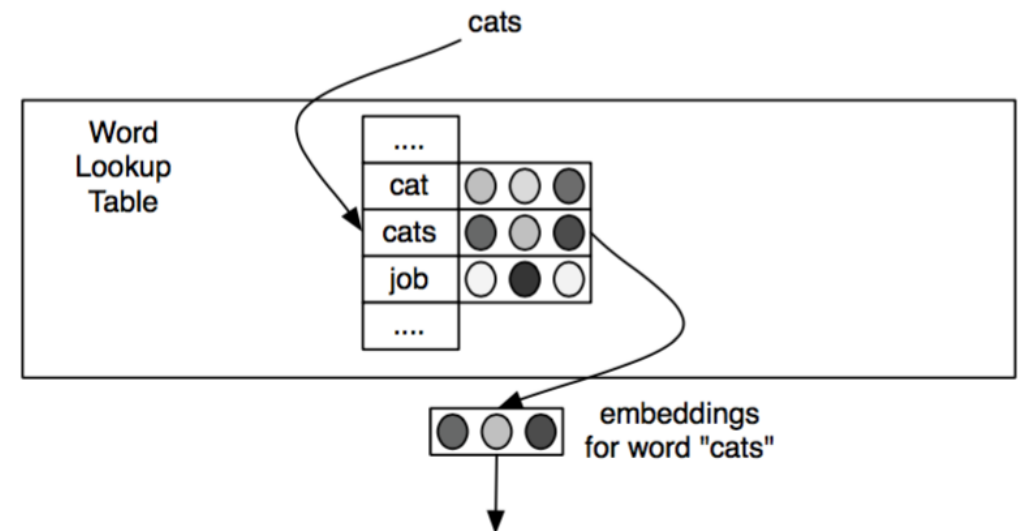
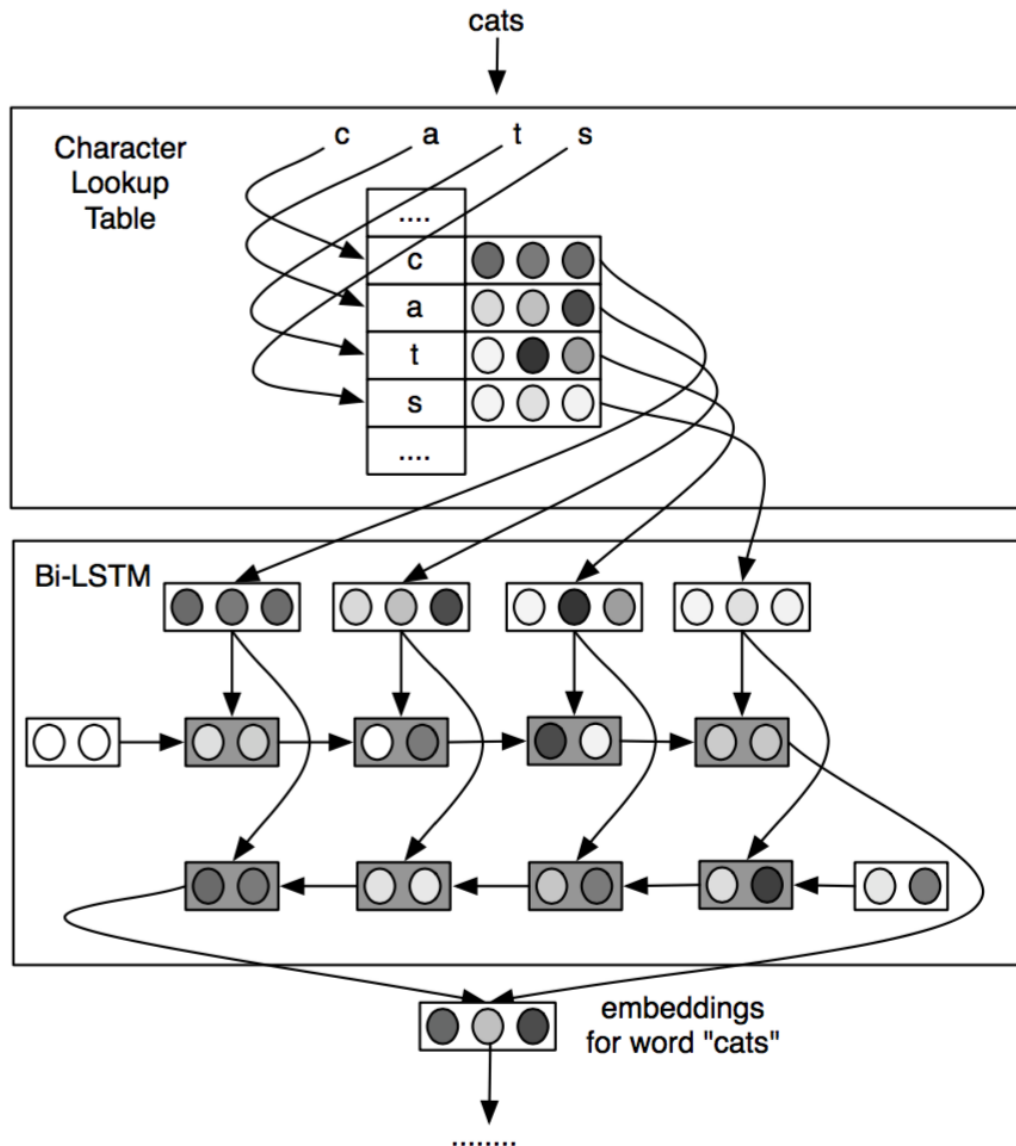


- Present the character sequence forwards and backwards to two separate LSTMs.
- Yields the forward state sequence  $s_1^f, \dots, s_m^f$  and backward state sequence  $s_m^b, \dots, s_1^b$ .
- The network has simultaneous access to all inputs before and after the current one.
- No need for fixed window sizes for the input, the net decides how much context to use.

# C2W-Model: Combining Layer

- ▶ Combines the last forward state  $s_m^f$  and the last backward state  $s_1^b$
- ▶  $e_w = D^f s_m^f + D^b s_0^b + b_d$
- ▶ The variables  $D^f, D^b, b_d$  are the weights which determine how the states are combined.
- ▶ Automatically learns to determine how much each context is used.

# C2W-Model vs Lookup Tables



# C2W-Model vs Word Lookup Tables

- ▶ **Lookup Tables** are conceptually much simpler, but requires a lot of parameters ( $|V| \times d$ ).
- ▶ The C2W model uses less parameters (more in the evaluation section).
- ▶ The C2W model can easily be used for open vocabulary tasks.
- ▶ Looking up a word embedding is in  $O(1)$ , whereas the C2W model has to compute the embedding
  - ▷ Can be alleviated by caching  $e_w$  for frequently occurring words.
  - ▷ However cached values still need to be recomputed when parameters change during training time.



# Experimentation

**We are going to introduce three use cases, where a C2W based model either:**

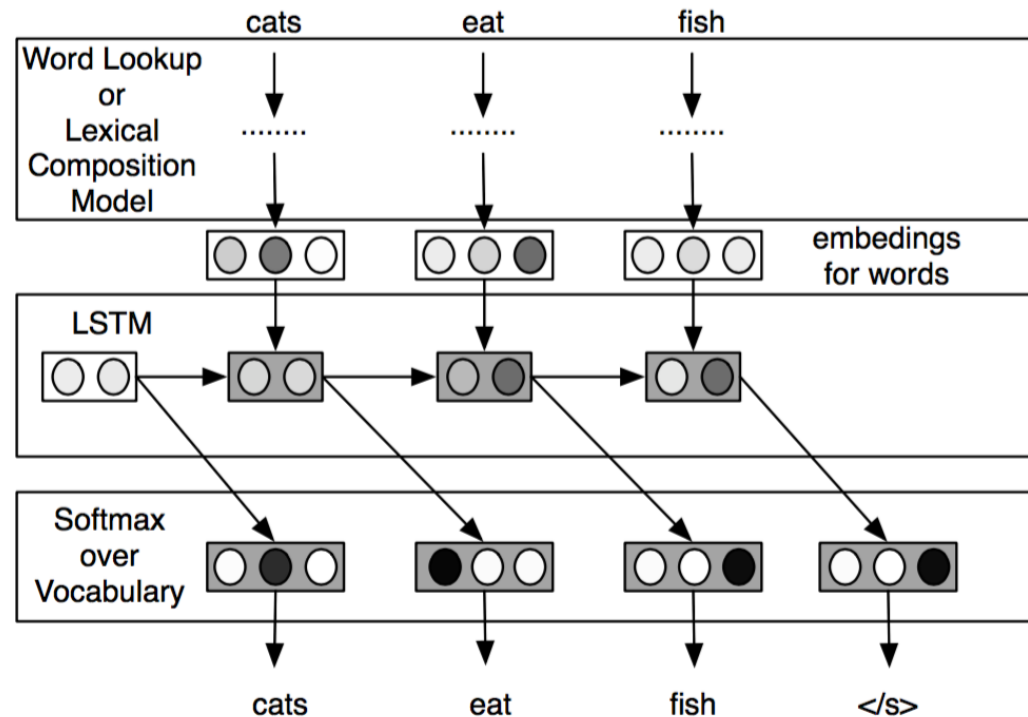
- ▶ **Outperforms a model which uses much more parameters.**
- ▶ **Yields comparable results without relying on manually engineered features.**

**Wang Ling et al 2015 : Language Modelling**

**Wang Ling et al 2015 : Part-Of-Speech Tagging**

**Faruqui et al 2015 : Morphological Inflection Generation**

# Application 1: Language Modeling



- ▶ Computes the joint probability for the word sequence  $w_1, \dots, w_{i-1}$  with a hidden LSTM layer.
- ▶ Test two versions of this NLM: One with the C2W model and one with a word lookup table as projection layer.
- ▶ Compare accuracy of these two versions.

# Application 1: Training & Testing Data

- ▶ **Perform testing on English, Portuguese, Catalan, German and Turkish.**
- ▶ **Chosen because these are morphologically rich languages.**
- ▶ **Training data was obtained by randomly extracting wikipedia articles until 1 million words were obtained.**
- ▶ **Additionally 20000 words were obtained for testing.**

# Application 1: Parameter Count

- ▶ The word embedding dimension is set to  $d = 50$ .
  - ▷ A word lookup table contains at least  $d \times |V|$  parameters,
  - ▷ A language with 80000 words will have at least 4 million parameters.
- ▶ The C2W model has two additional hyperparameters which are set to  $d_C = 50$  and  $d_{CS} = 150$ 
  - ▷ The LSTMs use 8 matrices of size  $d_{CS} \times d_C + 2d_{CS}$  (one for each decision gate).
  - ▷ The  $d \times 2d_{CS}$  parameters in the combining output layer.
  - ▷ The  $d_C \times |C|$  parameters in the character table.
  - ▷ For english this works out to roughly 180000 parameters.

# Reminder: Perplexity

- ▶ **Measure of how well a probability distribution predicts sample data.**
- ▶ **Can be interpreted as the number of choices per word position.**
- ▶ **Defined as  $2^{H(p)} = 2^{-\sum_x p(x) \log_2 p(x)}$**
- ▶ **To minimize the perplexity value means to have a better fitting language model.**

# Application 1: Evaluation

Perplexity	English	Portuguese	Catalan	German	Turkish
Word Lookup	59.38	46.17	35.34	43.02	44.01
C2W Model	<b>57.39</b>	<b>40.92</b>	<b>34.92</b>	<b>41.94</b>	<b>32.88</b>
#Parameters					
Word Lookup	4.3M	4.2M	4.3M	6.3M	5.7M
C2W Model	<b>180K</b>	<b>178K</b>	<b>182K</b>	<b>183K</b>	<b>174K</b>

Perplexities and test configuration [Wang Ling et al 2015].

- ▶ Training is performed with mini-batch gradient descent with 100 sentences each.
- ▶ Speed of both model versions is approximately 300 words per second, main bottleneck is the softmax layer.
- ▶ In general C2W outperforms word lookup tables and requires less parameters.

# Application 1: Nonce Words

Noahshire	phding
Nottinghamshire	mixing
Bucharest	modelling
Saxony	styling
Johannesburg	blaming
Gloucestershire	christening

**Nonce words and their most similar words from the vocabulary [Wang Ling et al 2015].**

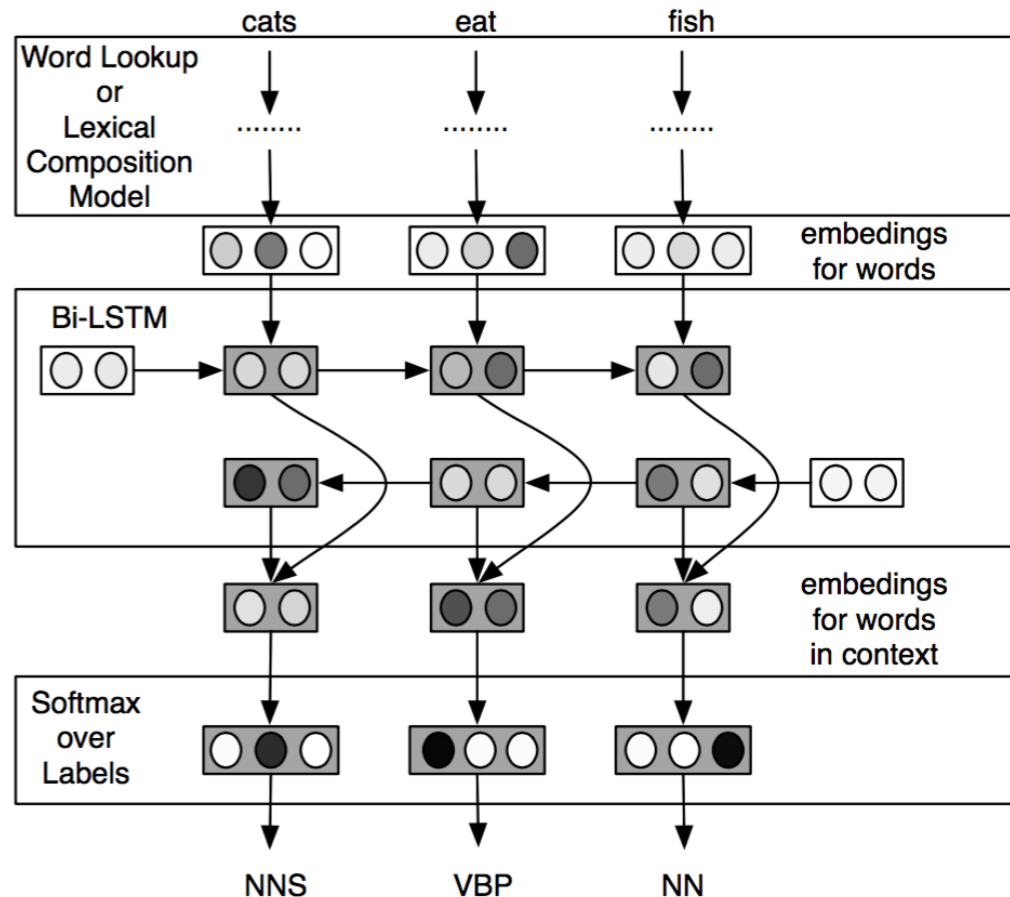
- ▶ **Nonce words are words created for use in a single occasion.**
- ▶ **The C2W model is able to generate embeddings for these words.**
- ▶ **No need for an OOV token for out of vocabulary words.**

## Application 2: Part-Of-Speech Tagging

- ▶ **Process of labeling words as corresponding to a particular part of speech**
- ▶ **For example a simple tagging would be to identify words as nouns, verbs, adjectives, adverbs, etc.**
- ▶ **One use of this is to disambiguate homonyms:  
For Example "I fish a fish" should become "Je pêche un poisson" in french.**
- ▶ **By tagging the first occurrence of "fish" as verb and the second as noun, they are now distinct.**



## Application 2: Part-Of-Speech Tagging



- ▶ Conceptually similar to the previous NLM model.
- ▶ We don't just use a single LSTM block, but a bidirectional LSTM.
- ▶ The softmax layer computes over all possible labels instead of the vocabulary

## Application 2: Testing Setup

- ▶ For english annotated sentences of the Wall Street Journal from the "Penn Treebank" dataset are used.
- ▶ For other languages data provided by the "Conference on Natural Language Learning" was used.
- ▶ The dimension of the states in the additional Bi-LSTM layer are set to 50.
- ▶ Out of vocabulary words are replaced with an OOV token.

## Application 2: Evaluation

	acc	parameters	words/sec
Word Lookup	96.97	2000k	6K
Convolutional (S&Z)	96.80	42.5k	4K
Forward RNN	95.66	17.5k	4K
Backward RNN	95.52	17.5k	4K
Bi-RNN	95.93	40k	3K
Forward LSTM	97.12	80k	3K
Backward LSTM	97.08	80k	3K
Bi-LSTM $d_{CS} = 50$	97.22	70k	3K
Bi-LSTM	<b>97.36</b>	150k	2K

- ▶ As previously the C2W based POS-model is compared with versions using different word representation models.
- ▶ Table contains accuracies for the english WSJ dataset only.
  - ▷ There are different configurations using regular RNNs and LSTMs.
  - ▷ The LSTMs always outperforms regular RNNs by about 2%.
  - ▷ Row "Convolutional (S&Z)" contains results of a convolutional model from [Santos and Zadrozny, 2014].

[Santos and Zadrozny, 2014]: Learning Character-level Representations for Part-of-Speech Tagging

## Application 2: Evaluation

System	Fusional			Agglutinative	
	EN	PT	CA	DE	TR
Word	96.97	95.67	98.09	97.51	83.43
C2W	<b>97.36</b>	97.47	<b>98.92</b>	<b>98.08</b>	<b>91.59</b>
Stanford	97.32	<b>97.54</b>	98.76	97.92	87.31

- ▶ This table contains testing results for a number of languages.
- ▶ As previously the row "word" contains results for a model version using word lookup tables.
- ▶ Additionally it is compared with Stanford's POS tagger, with the default set of features.

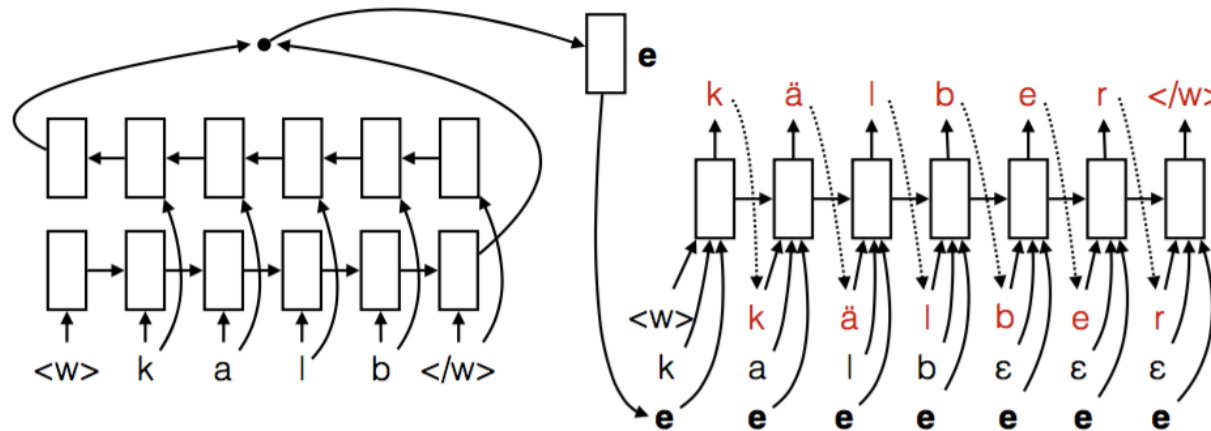
## Application 3: Morphological Inflection Generation

	singular	plural
nominative	Kalb	Kälber
accusative	Kalb	Kälber
dative	Kalb	Kälbern
genitive	Kalbes	Kälber

**Example of an inflection table for the word "Kalb" [Faruqui et al 2015]**


- ▶ **We want to perform morphological transformations of words.**
- ▶ **These kind transformations are very common in languages like turkish or german.**
- ▶ **Could be used as post- or preprocessing step for machine trainlation.**

## Application 3: Overview



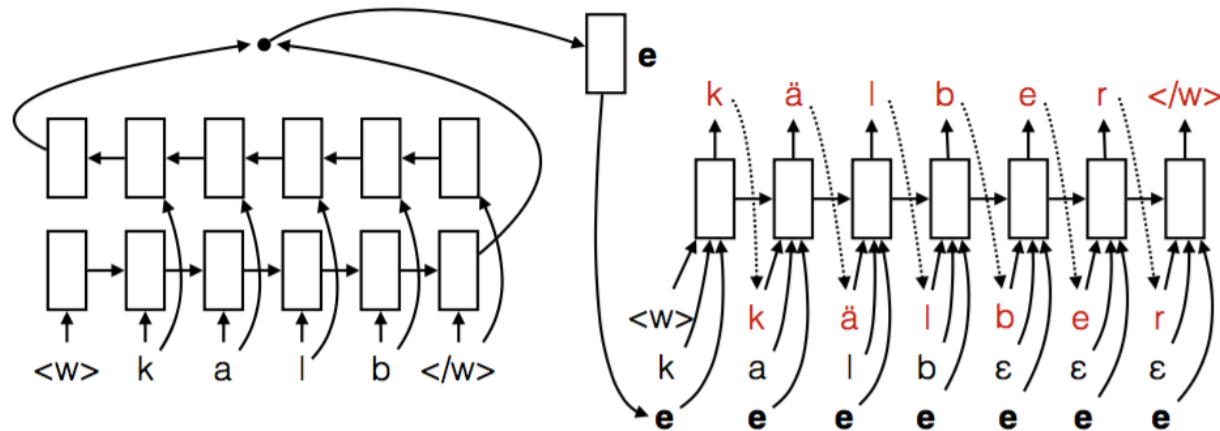
- ▶ We use a neuronal encoder - decoder architecture.
- ▶ The word embedding  $e$  is used as an intermediate representation.
- ▶ The decoder constructs the inflected version of the word - character by character.

# Application 3: Overview



- ▶ The encoder part is identically to the C2W model and generates a word embedding  $e$ .
- ▶ The decoder is just an LSTM unit which receives the following inputs each timestep:
  1. The word embedding  $e$  from the encoder.
  2. Current character of the original word  $c_j$
  3. Previous output of the model

# Application 3: Decoder



- ▶ The decoder output is driven by the characters of the original word
- ▶ There is a chance that output length is greater than the input word length.
- ▶ Once the input word ends, the  $\epsilon$  character is used instead.
- ▶ The decoder stops by outputting the word end token " $\langle /w \rangle$ "



## Application 3: Training & Testing Data

Dataset	root forms	Infl.
German Nouns (DE-N)	2764	8
German Verbs (DE-V)	2027	27
Spanish Verbs (ES-V)	4055	57
Finnish NN & Adj. (FI-NA)	6400	28
Finnish Verbs (FI-V)	7249	53
Dutch Verbs (NL-V)	11200	9
French Verbs (FR-V)	6957	48

The languages above were tested with the data published by:

- ▶ [Durrett and DeNero 2013] containing inflections for German, Finnish and Spanish.
- ▶ [Nicolai et al., 2015] adding dutch and french to this dataset.
- ▶ The development and test sets contain about 200 inflection tables each.

[Durrett and DeNero 2013]: Supervised learning of complete morphological paradigms. In Proc. of NAACL.

[Nicolai et al., 2015]: Inflection generation as discriminative string transduction. In Proc. of NAACL

## Application 3: Evaluation

	DDN13	NCK15	Ours
DE-V	94.76	<b>97.50</b>	96.72
DE-N	88.31	<b>88.60</b>	88.12
ES-V	99.61	99.80	<b>99.81</b>
FI-V	97.23	<b>98.10</b>	97.81
FI-NA	92.14	93.00	<b>95.44</b>
NL-V	90.50	96.10	<b>96.71</b>
FR-V	98.80	<b>99.20</b>	98.82
Avg.	94.47	96.04	<b>96.20</b>

- ▶ The results are comparable or better than other approaches.
- ▶ On average the results are better.
- ▶ No feature engineering necessary.

# Summary

- ▶ **Generating word embeddings by composing character representations, works usually just as well as approaches using word lookup tables.**
- ▶ **Lexical features can be learned automatically, manual feature engineering can be avoided**
- ▶ **In combination with caching of frequently used words, the performance is comparable to models based on word lookup tables.**
- ▶ **Models scale better with larger vocabularies and are able to deal with open vocabularies.**